

MPI_Gather: Gather together multiple values from different processors.

```
MPI_Gather (&input_value, 1, MPI_INT, &output_array, 1, MPI_INT, 0, MPI_COMM_WORLD)
```

APPROXIMATION OF π VIA MONTE CARLO SIMULATION

Draw a circle inside of a square and randomly place dots in the square. The ratio of dots inside the circle to the total number of dots will approximately equal $\pi/4$.

```
// Sequential
long count_hits(long trials) { long hits = 0, i; for (i = 0; i < trials; i++) {
    double x = (double)rand() / RAND_MAX; double y = (double)rand() / RAND_MAX;
    if (x * x + y * y <= 1) { hits++; } // distance to center bigger than radius=1
} return hits; }

// Parallel, the trials are split across different nodes
int rank, size;
MPI_Comm_rank(MPI_COMM_WORLD, &rank); MPI_Comm_size(MPI_COMM_WORLD, &size);
srand(rank * 4711); // each process receives a different seed
long hits = count_hits(TRIALS / size); // Each process computes a subtask
long total;
MPI_Reduce(&hits, &total, 1, MPI_LONG, MPI_SUM, 0, MPI_COMM_WORLD);
if (rank == 0) { double pi = 4 * ((double)total / TRIALS); }
```

OPENMP

Node: A standalone “computer in a box”. Usually comprised of multiple CPU/processors/cores, memory, network interfaces etc. Nodes are **networked together** to comprise a **supercomputer**. Each node consists of 20 cores. The processes do **not share memory**, they must use messages. **Threads:** Default are 24 on a single Node in OST cluster. Can be set with `omp_set_num_threads()` or with the `OMP_NUM_THREADS` environment variable. Threads range from 0 (*master thread*) to N-1. **HPC Hybrid memory model:** Run a program on multiple nodes. No Shared Memory (*NUMA*) between Nodes. Shared Memory (*SMP*) for Cores inside a Node.

OpenMP: Is a programming model for different languages. **Allows to run multiple threads**, distribute work using synchronization and reduction constructs. **Shared Memory** (*shared memory process consists of multiple threads*), **Explicit Parallelism** (*Programmer has full control over parallelization*) and **Compiler Directives** (*Most OpenMP parallelism is specified through the use of compiler directives (pragmas) in the source code*).

```
Fork and Join
#include <stdio.h>
#include <omp.h>
int main(int argc, char* argv[]) {
    const int np = omp_get_max_threads(); // executed by initial thread
    printf("OpenMP with threads %d\n", np); // executed by initial thread
#pragma omp parallel // pragma spans multiple threads (fork)
    {
        const int np = omp_get_num_threads(); // executed in parallel
        printf("Hello from thread %d\n", omp_get_thread_num()); // executed in paral.
    } // thread order not fixed. After execution, threads synchronize & terminate
    return 0; }
```

For loops

```
#pragma omp parallel for
for (i=0; i<n; i++) { ... }
```

Each thread processes **one loop-iteration** at a time. Execution returns to the initial threads.

Oversubscription (*too many threads for a problem*) is handled by OpenMP. The iteration variable (*i.e. j*) is implicitly made private for the duration of the loop.

Memory Model

```
int A, B, C // automatically global because outside of pragma
#pragma omp parallel for private (A) shared (B) firstprivate (C)
for(...)
```

Each thread has a **private copy of A** and use the **same memory location for B**. C is also private, but gets its initial value from the global variable. After the loop is over, threads die and both A and B will be cleared/removed from memory.

```
#pragma omp parallel
    int A = 0 // automatically private because inside of pragma
#pragma omp for ...
```

Avoiding Race conditions: Mutex

```
int sum = 0;
#pragma omp parallel for
for (int i = 0; i < n; i++)
#pragma omp critical { sum += i; } // only one thread at a time
```

This is **extremely slow** due to serialization, slower than single threading. Critical section is **overkill** for this code, with a heavy weight mutex the performance overhead is large.

Lightweight mutex: Atomic

```
int sum = 0; int i;
#pragma omp parallel for
for (i = 0; i < n; i++)
    #pragma omp atomic { sum += i }
```

Reduction across threads

When using reduction (operator: variable), a **copy** of the reduction variable per thread is created, initialized to the identity of the reduction operator ($+=$ 0 , $+= 1$). Each thread will then **reduce** into its local variable. At the end of the parallel region, the local results are **combined into the global variable**. Only associative operators allowed ($+$, $*$, not , $-$, $>$).

```
// Code using the reduction clause
int sum = 0;
#pragma omp parallel for reduction(+: sum)
for (int i = 0; i < n; i++) { sum += i; }
```

```
// The same code without the reduction clause
int sum = 0;
#pragma omp parallel {
    int intermediate_sum = 0; // private
    #pragma omp for
    for (int i = 0; i < n; i++) { intermediate_sum += i; } // thread partial sum
#pragma omp atomic // reduction is protected with atomic
    final_sum += intermediate_sum; }
```

Hybrid: OpenMP + MPI

```
int numprocs, rank; int iam = 0, np = 1;
MPI_Init(&argc, &argv);
MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
#pragma omp parallel default(shared) private(iam, np) {
    np = omp_get_num_threads();
    iam = omp_get_thread_num();
    printf("I am %d out of %d from P %d out of %d\n", iam, np, rank, numprocs);
} MPI_Finalize();
```

Sequential count_hits to approximate π with Monte Carlo Simulation in OpenMP

```
long count_hits(long trials) {
    long hits = 0; long i; double x,y;
#pragma omp parallel {
    #pragma omp for reduction(+:hits) private(x,y)
    for (i = 0; i < trials; i++) {
        double x = random_double(); double y = random_double();
        if (x * x + y * y <= 1) { hits += 1; }
    }
} return hits;
}
```

PERFORMANCE SCALING

Difficulties with parallel programs: Finding parallelism, granularity of a parallel task, moving data is expensive, load balancing, coordination & synchronization, performance debugging.

Scalability: The ability of hard- and software to deliver **greater computational power** when the number of **resources is increased**.

Scalability Testing: Primary challenge of parallel computing is **deciding how best to break up a problem** into individual pieces that can be computed separately. It is impractical to develop and test large applications using the **full problem size**. The problem and number of processors are **scaled down** at first. **Scalability testing:** measuring the ability of an application to perform well or better with varying problem sizes and numbers of processors. It does **not test** the applications **general functionality** or correctness.

STRONG SCALING

The number of processors is increased while the problem size remains constant. Results in a reduced workload per processor. Mostly used for long running CPU bound applications.

Amdahls Law: The speedup is **limited by the fraction of the serial part** of the software that is not amenable to parallelization. Sweet spot needs to be found. It is reasonable to use **small amounts of resources for small problems** and **larger quantities of resources for big problems**.

$T =$ total time, $p =$ part of the program that can be parallelized, $N =$ amount of processors
 $T = (1 - p)T + T_p + T_p$
 $T_N = T_p / N + (1 - p)T_p$, Speedup $\leq 1 / (s + p / N)$, Efficiency $= T / (N \cdot T_N)$
Amdahls law **ignores the parallel overhead**. Because of that, it is the **upper limit** of speedup for a problem of fixed size. This seems to be a **bottleneck** for parallel computing.
Examples: 90% of the computation can run parallel, what is the max speedup with 8 processors?
 $1 / (0.1 + 0.9 / 8) \approx 4.7$
25% of the computation must be serial. What is the max speedup with ∞ Processors?
 $1 / (0.25 + 0.75 / \infty) \approx 1 / 0.25 = 4$

To gain a 500 x speedup on 1000 processors, Amdahls law requires that the proportion of serial part cannot exceed what?
 $500 = 1 / (s + \frac{1-p}{1000}) \Rightarrow s + \frac{1-p}{1000} = \frac{1}{500} \Rightarrow 1000s + (1 - s) = 2 \Rightarrow 999s = 1 \Rightarrow s = \frac{1}{999} \approx 0.1\%$

WEAK SCALING

The number of processors and the problem size is increased. Mostly used for large **memory-bound** applications where the required memory cannot be satisfied by a single node.

Gustafson's law: Based on the approximations that the **parallel part scales linearly** with the amount of resources, and that the **serial part** does **not increase** with respect to the size of the problem. **Speedup** $= s + p \cdot N = s + (1 - s) \cdot N$

In this case, the problem size assigned to each processing element stays **constant** and **additional elements** are used to solve a **larger total problem**. Therefore, this type of measurement is justification for programs that take a lot of memory or other system resources.

Example: 64 Processors, 5% of the program is serial, What is the scaled weak speedup?

$0.05 + 0.95 \cdot 64 = 60.85$